

stackedGen

November 13, 2019

1 Stacking

Stacked generalization is an ensemble technique that involves two stages. The first stage trains a set of M models on N observations and stores their predictions in a $L_1 = N \times M$ data set. The second stage uses the dataset of predictions to train an ensemble model, called a **generalizer**. The generalizer is meant to learn the strengths and weaknesses of each model through their predictions, and built a more accurate model.

The goal of a stacked generalization algorithm is not to better understand the relationship between a set of covariates X and a target variable y . The goal is to make as accurate a model as possible.

1.1 Stages

We will use the Boston housing market dataset to understand stacking. The goal will be to predict the median value of homes in Boston given several characteristics of the housing market.

```
[1]: require(mlbench)
data(BostonHousing)

dim(BostonHousing)
summary(BostonHousing)
```

Loading required package: mlbench

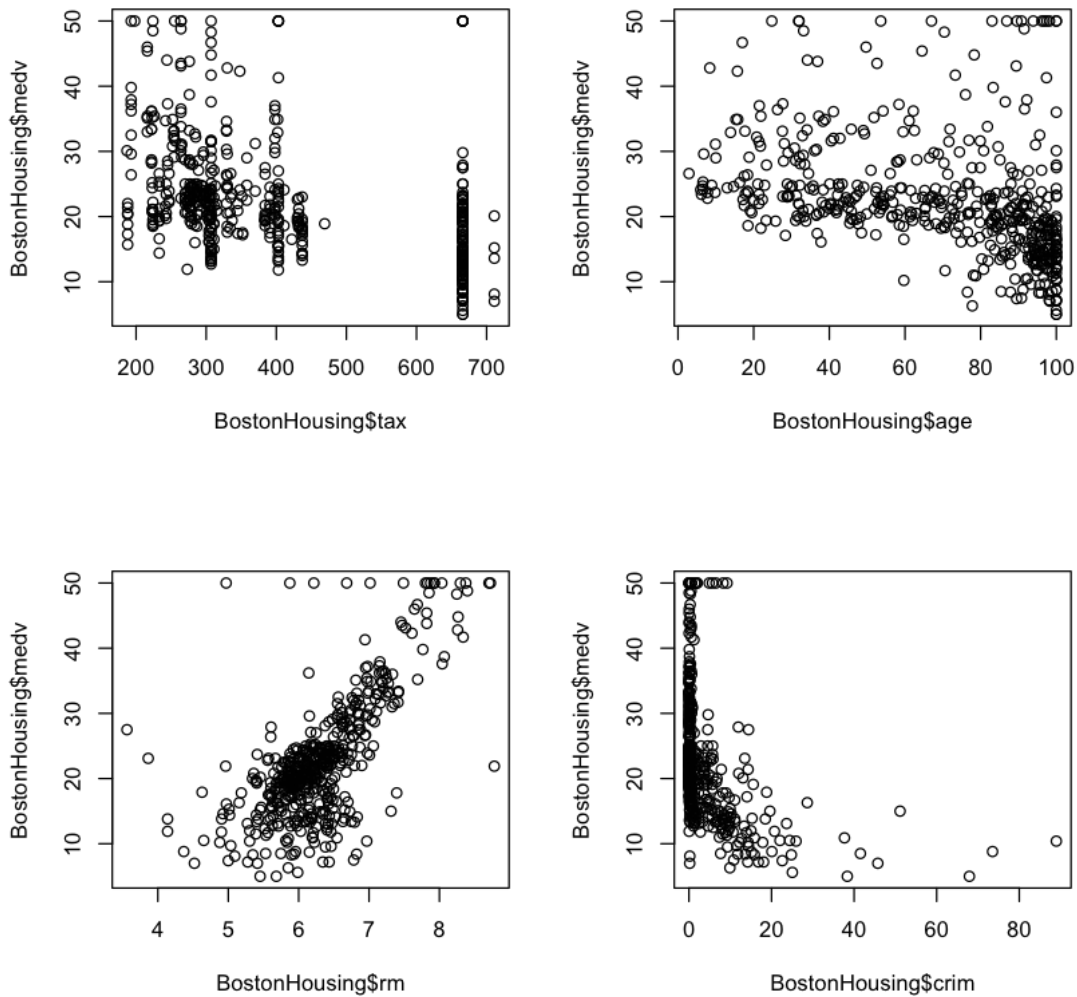
1.506 2.14

crim	zn	indus	chas	nox
Min. : 0.00632	Min. : 0.00	Min. : 0.46	0:471	Min. :0.3850
1st Qu.: 0.08204	1st Qu.: 0.00	1st Qu.: 5.19	1: 35	1st Qu.:0.4490
Median : 0.25651	Median : 0.00	Median : 9.69		Median :0.5380
Mean : 3.61352	Mean : 11.36	Mean :11.14		Mean :0.5547
3rd Qu.: 3.67708	3rd Qu.: 12.50	3rd Qu.:18.10		3rd Qu.:0.6240
Max. :88.97620	Max. :100.00	Max. :27.74		Max. :0.8710

rm	age	dis	rad
Min. :3.561	Min. : 2.90	Min. : 1.130	Min. : 1.000
1st Qu.:5.886	1st Qu.: 45.02	1st Qu.: 2.100	1st Qu.: 4.000
Median :6.208	Median : 77.50	Median : 3.207	Median : 5.000

Mean	:6.285	Mean	: 68.57	Mean	: 3.795	Mean	: 9.549
3rd Qu.	:6.623	3rd Qu.	: 94.08	3rd Qu.	: 5.188	3rd Qu.	:24.000
Max.	:8.780	Max.	:100.00	Max.	:12.127	Max.	:24.000
tax		ptratio		b		lstat	
Min.	:187.0	Min.	:12.60	Min.	: 0.32	Min.	: 1.73
1st Qu.	:279.0	1st Qu.	:17.40	1st Qu.	:375.38	1st Qu.	: 6.95
Median	:330.0	Median	:19.05	Median	:391.44	Median	:11.36
Mean	:408.2	Mean	:18.46	Mean	:356.67	Mean	:12.65
3rd Qu.	:666.0	3rd Qu.	:20.20	3rd Qu.	:396.23	3rd Qu.	:16.95
Max.	:711.0	Max.	:22.00	Max.	:396.90	Max.	:37.97
medv							
Min.	: 5.00						
1st Qu.	:17.02						
Median	:21.20						
Mean	:22.53						
3rd Qu.	:25.00						
Max.	:50.00						

```
[2]: par(mfrow=c(2,2))
plot(BostonHousing$tax,BostonHousing$medv)
plot(BostonHousing$age,BostonHousing$medv)
plot(BostonHousing$rm,BostonHousing$medv)
plot(BostonHousing$crim,BostonHousing$medv)
```



1.2 (1) Partition data into training and testing.

Stacking occurs in stages. The first stage partitions your data into a training and test set.

```
[3]: percentTraining = 0.80
BostonHousing['train'] = runif(nrow(BostonHousing)) < percentTraining

training = BostonHousing[BostonHousing$train==1,]
testing  = BostonHousing[BostonHousing$train==0,]

print(head(BostonHousing,10))
```

```
print('size of training data')
dim(training)

print('size of testing data')
dim(testing)
```

```
      crim   zn indus chas   nox    rm   age   dis rad tax ptratio    b
1  0.00632 18.0  2.31    0 0.538 6.575  65.2 4.0900   1 296    15.3 396.90
2  0.02731  0.0  7.07    0 0.469 6.421  78.9 4.9671   2 242    17.8 396.90
3  0.02729  0.0  7.07    0 0.469 7.185  61.1 4.9671   2 242    17.8 392.83
4  0.03237  0.0  2.18    0 0.458 6.998  45.8 6.0622   3 222    18.7 394.63
5  0.06905  0.0  2.18    0 0.458 7.147  54.2 6.0622   3 222    18.7 396.90
6  0.02985  0.0  2.18    0 0.458 6.430  58.7 6.0622   3 222    18.7 394.12
7  0.08829 12.5  7.87    0 0.524 6.012  66.6 5.5605   5 311    15.2 395.60
8  0.14455 12.5  7.87    0 0.524 6.172  96.1 5.9505   5 311    15.2 396.90
9  0.21124 12.5  7.87    0 0.524 5.631 100.0 6.0821   5 311    15.2 386.63
10 0.17004 12.5  7.87    0 0.524 6.004  85.9 6.5921   5 311    15.2 386.71
  lstat medv train
1   4.98 24.0  TRUE
2   9.14 21.6  TRUE
3   4.03 34.7  TRUE
4   2.94 33.4  TRUE
5   5.33 36.2  TRUE
6   5.21 28.7  TRUE
7  12.43 22.9  TRUE
8  19.15 27.1  TRUE
9  29.93 16.5  TRUE
10 17.10 18.9  TRUE
[1] "size of training data"

1.410 2.15

[1] "size of testing data"

1.96 2.15
```

1.3 (2) Cross validation for out-of-sample predictions

Next we split our training data into K folds. For every fold k , we train on the left over $K - 1$ folds and make predictions on fold k . We repeat this process for all M component models.

We will use a KNN neighbor regression, linear regression, polynomial regression, and regression tree to predict the median value of houses (medv).

```
[4]: require(FNN) # for the KNN model
      require(rpart) # for the TBR model

K = 10
```

```

training = training[,names(training)!='train']
trainingFolds = split(training,1:K)

dataSetOfPredictions = matrix()
for(k in 1:K){
  sprintf("Fold %d", k)

  outOfSample = trainingFolds[[k]]
  outOfSampleX = outOfSample[,names(outOfSample)!='medv']

  leftOver = setdiff(1:K,k)
  trainingSamples = do.call(rbind,trainingFolds[leftOver])

  trainY = trainingSamples[,names(trainingSamples)=='medv']
  trainX = trainingSamples[,names(trainingSamples)!='medv']

  # KNN model
  m1_predictions = knn.reg(train = as.matrix(trainingSamples$crim)
    ,test = as.matrix(outOfSampleX$crim)
    ,y=as.matrix(trainY),k=10)$pred

  #linear regression
  m2 = lm(medv~., data = trainingSamples)
  m2_predictions = predict(m2,outOfSampleX)

  #polynomial regression
  m3 = lm(medv~tax + age + rm + crim + I(tax^2) + I(age^2) + I(rm^2) +
→I(crim^2), data = trainingSamples)
  m3_predictions = predict(m3,outOfSampleX)

  #TBR
  m4 = rpart(medv ~ ., method="anova", data=trainingSamples)
  m4_predictions = predict(m4,outOfSampleX)

  # build data set of out-of-sample predictions
  if (k==1){
    allDataSetOfPredictions =
→cbind(m1_predictions,m2_predictions,m3_predictions,m4_predictions,outOfSample$medv)
  } else {
    dataSetOfPredictions =
→cbind(m1_predictions,m2_predictions,m3_predictions,m4_predictions,outOfSample$medv)
    allDataSetOfPredictions =
→rbind(allDataSetOfPredictions,dataSetOfPredictions,outOfSample$medv)
  }
}

```

Loading required package: FNN

Loading required package: rpart

```
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions,
outOfSample$medv):
â€”Number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions,
outOfSample$medv):
â€”Number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions,
outOfSample$medv):
â€”Number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions,
outOfSample$medv):
â€”Number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions,
outOfSample$medv):
â€”Number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions,
outOfSample$medv):
â€”Number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions,
outOfSample$medv):
â€”Number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions,
outOfSample$medv):
â€”Number of columns of result is not a multiple of vector length (arg 3)â€”
```

```
[5]: cbind(m1_predictions,m2_predictions,m3_predictions,m4_predictions,outOfSample$medv)
```

	m1_predictions	m2_predictions	m3_predictions	m4_predictions	
10	21.46	19.027995	20.009662	18.359615	18.9
22	20.87	17.810408	19.477677	20.613793	19.6
34	16.27	14.443087	18.148213	18.359615	13.1
47	21.18	20.583063	22.681509	20.613793	20.0
62	21.46	18.960698	19.837244	18.359615	16.0
75	27.02	25.537439	24.221358	22.750877	24.1
85	23.64	24.591988	24.919208	22.750877	23.9
99	27.32	33.994662	40.953380	38.985714	43.8
111	21.85	20.466214	21.705076	20.613793	21.7
123	23.23	21.456362	21.415900	18.359615	20.5
136	36.97	17.452666	19.949913	14.569231	18.1
148	18.52	9.053367	16.316344	14.569231	14.6
158	16.25	33.179382	25.085168	27.868293	41.3
172	21.91	24.644861	17.224248	20.613793	19.1
187	24.94	34.923693	41.720271	38.985714	50.0
200	26.25	30.213712	29.624447	27.868293	34.9
213	19.14	22.254644	21.145573	18.359615	22.4
224	29.22	29.513213	24.083355	27.868293	30.1
237	32.22	29.175402	24.427489	27.868293	25.1
249	21.90	21.472451	24.011281	22.750877	24.5
259	27.13	35.801714	31.175662	32.234615	36.0
273	24.27	28.610831	25.888060	22.750877	24.4
291	27.41	33.957245	29.779209	27.868293	28.5
306	24.77	30.769238	26.579860	27.868293	28.4
317	23.76	17.670191	19.780897	18.359615	17.8
331	22.60	21.167586	22.090658	22.750877	19.8
341	22.56	21.728596	22.432882	22.750877	18.7
352	27.02	20.792584	24.974866	27.868293	24.1
362	19.63	19.196094	18.421405	20.613793	19.9
373	15.13	25.780359	15.137599	33.057143	50.0
387	10.47	6.670162	9.923174	9.361538	10.5
398	14.62	16.625857	14.448075	14.569231	8.5
409	14.79	13.831855	14.236169	18.359615	17.2
420	13.57	14.164846	20.958169	9.361538	8.4
434	22.50	17.084606	19.196868	14.569231	14.3
448	16.94	18.255262	16.135249	14.569231	12.6
460	17.96	18.817369	16.780565	14.569231	20.0
470	13.94	18.949890	14.634050	18.359615	20.1
483	22.38	28.447699	25.308605	32.234615	25.0
495	21.70	20.586318	20.789736	20.613793	24.5
506	23.36	22.706229	20.925367	33.057143	11.9

We now have a dataset that includes an out-of-sample for prediction, for all 4 models and for all observation in our training set.

```
[6]: allDataSetOfPredictions = data.frame(allDataSetOfPredictions)
names(allDataSetOfPredictions) = c('m1', 'm2', 'm3', 'm4', 'y')
```

```
print(head(allDataSetOfPredictions))
```

```

      m1      m2      m3      m4      y
X1  32.47 30.30560 24.90512 23.60122 24.0
X11 20.21 18.50869 21.81050 16.97647 15.0
X23 18.72 15.52379 20.15813 16.97018 15.2
X36 24.24 24.12824 21.02700 23.60122 18.9
X48 20.78 17.68533 21.36763 16.97018 16.6
X63 21.35 24.29733 23.99810 23.60122 22.2

```

1.4 (3) Build Aggregator

The next step trains a model that aggregates the M models together by training on the data set of out-of-sample predictions. We can consider a linear regression model as our aggregator.

```
[7]: agg = lm(y~m1+m2+m3+m4,data = data.frame(allDataSetOfPredictions))
      print(summary(agg))
```

Call:

```
lm(formula = y ~ m1 + m2 + m3 + m4, data = data.frame(allDataSetOfPredictions))
```

Residuals:

```

      Min       1Q   Median       3Q      Max
-25.1627  -2.2186  -0.3911   1.8806  28.4699

```

Coefficients:

```

              Estimate Std. Error t value Pr(>|t|)
(Intercept) -1.25579    0.94770  -1.325   0.1859
m1           -0.04067    0.04703  -0.865   0.3877
m2            0.48766    0.05554   8.780 <2e-16 ***
m3            0.14123    0.05736   2.462  0.0142 *
m4            0.47185    0.05293   8.915 <2e-16 ***

```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 4.408 on 414 degrees of freedom

Multiple R-squared: 0.7628, Adjusted R-squared: 0.7605

F-statistic: 332.8 on 4 and 414 DF, p-value: < 2.2e-16

These estimates are for our training set. We still need to make predictions on our held out test set.

1.5 (4) make predictions on test set

1.5.1 (4.1) Train Models on whole training set

```
[8]: # KNN model
      m1_predictions = knn.reg(train = as.matrix(training$crim)
                               ,test = as.matrix(testing$crim)
                               ,y=as.matrix(training$medv),k=10)$pred

      #linear regression
      m2 = lm(medv~., data = training)
      m2_predictions = predict(m2,testing)

      #polynomial regression
      m3 = lm(medv~tax + age + rm + crim + I(tax^2) + I(age^2) + I(rm^2) +
      ↪I(crim^2), data = training)
      m3_predictions = predict(m3,testing)

      #TBR
      m4 = rpart(medv ~ ., method="anova", data=training)
      m4_predictions = predict(m4,testing)

testSetPredictions =
  ↪cbind(m1_predictions,m2_predictions,m3_predictions,m4_predictions)
testSetPredictions = data.frame(testSetPredictions)
names(testSetPredictions) = c('m1','m2','m3','m4')

AggregatorModelPredictionsOnTestSet = predict(agg,data.frame(testSetPredictions))
print(AggregatorModelPredictionsOnTestSet)
```

	20	21	26	30	35	40	45	52
19.425249	14.770881	15.017293	21.199896	15.189805	28.215449	23.499058	23.922685	
53	56	58	60	64	71	73	87	
26.177563	32.121589	32.353396	22.429286	27.197722	25.004058	24.485570	21.572056	
88	90	93	105	108	115	120	129	
24.582588	31.364341	26.688665	20.814034	20.544854	22.750873	20.511577	17.991880	
130	131	137	141	159	160	161	163	
15.363997	20.353915	16.265534	12.211035	35.172640	33.866898	37.488204	43.554484	
177	180	181	183	186	189	190	194	
23.059418	32.390298	40.853508	32.909490	22.865273	28.476424	33.447676	32.102800	
201	206	211	218	225	233	234	240	
31.554960	22.049231	19.629651	24.650720	43.769164	43.453029	43.218009	26.517492	
248	261	263	270	271	276	278	279	
20.425024	33.048085	45.073319	23.197886	21.712013	32.983478	33.690091	27.347859	
280	282	283	285	288	295	296	298	
33.857480	33.187698	43.987926	31.997501	25.514842	22.666929	26.841570	18.334608	

	301	305	314	322	324	325	330	346
	31.390072	32.980208	24.685590	24.647241	19.933483	24.752430	24.088464	18.529734
	364	379	383	384	386	394	400	413
	20.151830	12.897429	11.719090	11.563689	8.823183	18.208772	10.571225	8.947883
	423	424	425	426	436	437	441	444
	19.098862	11.735513	15.250283	9.736015	12.219236	15.431773	11.269219	17.527636
	451	455	474	477	478	487	489	505
	16.851743	16.132665	23.666398	18.608707	10.621200	18.068553	14.390647	28.811103

1.6 (5) Compare your aggregated model to the test set

```
[9]: testSetAndAggPredictions = data.frame('T' = testing$medv
                                           , 'P' = AggregatorModelPredictionsOnTestSet
                                           , 'M1' = m1_predictions
                                           , 'M2' = m2_predictions
                                           , 'M3' = m3_predictions
                                           , 'M4' = m4_predictions)

SSE_AGG = sum( (testSetAndAggPredictions$T - testSetAndAggPredictions$P)^2 )

SSE_M1 = sum( (testSetAndAggPredictions$T - testSetAndAggPredictions$M1)^2 )
SSE_M2 = sum( (testSetAndAggPredictions$T - testSetAndAggPredictions$M2)^2 )
SSE_M3 = sum( (testSetAndAggPredictions$T - testSetAndAggPredictions$M3)^2 )
SSE_M4 = sum( (testSetAndAggPredictions$T - testSetAndAggPredictions$M4)^2 )


print(SSE_AGG)
print(SSE_M1)
print(SSE_M2)
print(SSE_M3)
print(SSE_M4)
```

```
[1] 1076.369
[1] 5304.312
[1] 1440.279
[1] 1380.285
[1] 2250.818
```

```
[22]: plot(testing$age, testing$medv)

S = order(testing$age)

lines(testing$age[S], testSetAndAggPredictions$P[S], col='red')
```



[]: