

stackedGen

November 12, 2019

1 Stacking

Stacked generalization is an ensemble technique that involves two stages. The first stage trains a set of M models on N observations and stores their predictions in a $L_1 = N \times M$ data set. The second stage uses the dataset of predictions to train an ensemble model, called a **generalizer**. The generalizer is meant to learn the strengths and weaknesses of each model through their predictions, and built a more accurate model.

The goal of a stacked generalization algorithm is not to better understand the relationship between a set of covariates X and a target variable y . The goal is to make as accurate a model as possible.

1.1 Stages

We will use the Boston housing market dataset to understand stacking. The goal will be to predict the median value of homes in Boston given several characteristics of the housing market.

```
[1]: require(mlbench)
data(BostonHousing)

dim(BostonHousing)
summary(BostonHousing)
```

Loading required package: mlbench

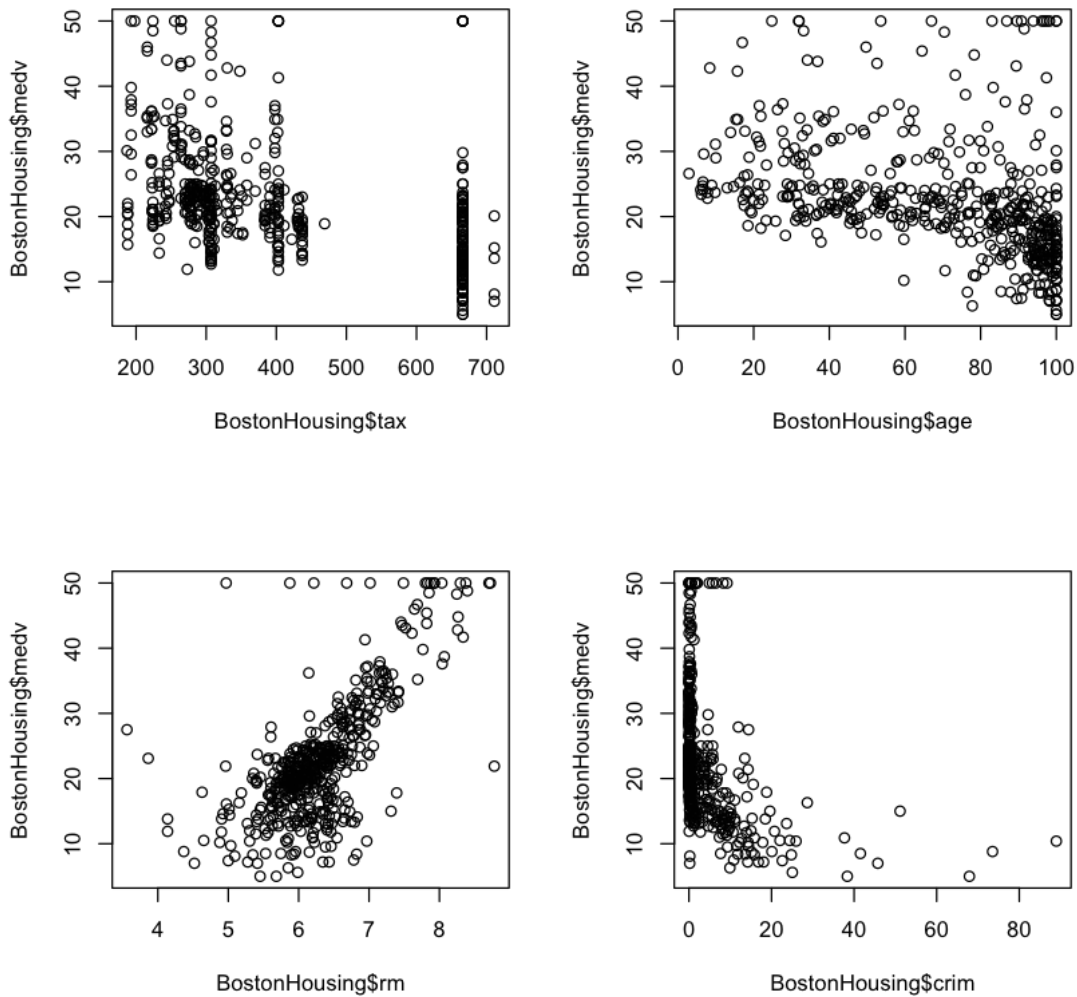
1.506 2.14

crim	zn	indus	chas	nox
Min. : 0.00632	Min. : 0.00	Min. : 0.46	0:471	Min. :0.3850
1st Qu.: 0.08204	1st Qu.: 0.00	1st Qu.: 5.19	1: 35	1st Qu.:0.4490
Median : 0.25651	Median : 0.00	Median : 9.69		Median :0.5380
Mean : 3.61352	Mean : 11.36	Mean :11.14		Mean :0.5547
3rd Qu.: 3.67708	3rd Qu.: 12.50	3rd Qu.:18.10		3rd Qu.:0.6240
Max. :88.97620	Max. :100.00	Max. :27.74		Max. :0.8710

rm	age	dis	rad
Min. :3.561	Min. : 2.90	Min. : 1.130	Min. : 1.000
1st Qu.:5.886	1st Qu.: 45.02	1st Qu.: 2.100	1st Qu.: 4.000
Median :6.208	Median : 77.50	Median : 3.207	Median : 5.000

Mean	:6.285	Mean	: 68.57	Mean	: 3.795	Mean	: 9.549
3rd Qu.	:6.623	3rd Qu.	: 94.08	3rd Qu.	: 5.188	3rd Qu.	:24.000
Max.	:8.780	Max.	:100.00	Max.	:12.127	Max.	:24.000
tax		ptratio		b		lstat	
Min.	:187.0	Min.	:12.60	Min.	: 0.32	Min.	: 1.73
1st Qu.	:279.0	1st Qu.	:17.40	1st Qu.	:375.38	1st Qu.	: 6.95
Median	:330.0	Median	:19.05	Median	:391.44	Median	:11.36
Mean	:408.2	Mean	:18.46	Mean	:356.67	Mean	:12.65
3rd Qu.	:666.0	3rd Qu.	:20.20	3rd Qu.	:396.23	3rd Qu.	:16.95
Max.	:711.0	Max.	:22.00	Max.	:396.90	Max.	:37.97
medv							
Min.	: 5.00						
1st Qu.	:17.02						
Median	:21.20						
Mean	:22.53						
3rd Qu.	:25.00						
Max.	:50.00						

```
[2]: par(mfrow=c(2,2))
plot(BostonHousing$tax,BostonHousing$medv)
plot(BostonHousing$age,BostonHousing$medv)
plot(BostonHousing$rm,BostonHousing$medv)
plot(BostonHousing$crim,BostonHousing$medv)
```



1.2 (1) Partition data into training and testing.

Stacking occurs in stages. The first stage partitions your data into a training and test set.

```
[33]: percentTraining = 0.80
BostonHousing['train'] = runif(nrow(BostonHousing)) < percentTraining

training = BostonHousing[BostonHousing$train==1,]
testing  = BostonHousing[BostonHousing$train==0,]

print(head(BostonHousing,10))
```

```
print('size of training data')
dim(training)

print('size of testing data')
dim(testing)
```

```
      crim   zn indus chas   nox    rm   age   dis rad tax ptratio    b
1  0.00632 18.0  2.31    0 0.538 6.575  65.2 4.0900   1 296    15.3 396.90
2  0.02731  0.0  7.07    0 0.469 6.421  78.9 4.9671   2 242    17.8 396.90
3  0.02729  0.0  7.07    0 0.469 7.185  61.1 4.9671   2 242    17.8 392.83
4  0.03237  0.0  2.18    0 0.458 6.998  45.8 6.0622   3 222    18.7 394.63
5  0.06905  0.0  2.18    0 0.458 7.147  54.2 6.0622   3 222    18.7 396.90
6  0.02985  0.0  2.18    0 0.458 6.430  58.7 6.0622   3 222    18.7 394.12
7  0.08829 12.5  7.87    0 0.524 6.012  66.6 5.5605   5 311    15.2 395.60
8  0.14455 12.5  7.87    0 0.524 6.172  96.1 5.9505   5 311    15.2 396.90
9  0.21124 12.5  7.87    0 0.524 5.631 100.0 6.0821   5 311    15.2 386.63
10 0.17004 12.5  7.87    0 0.524 6.004  85.9 6.5921   5 311    15.2 386.71
  lstat medv train
1   4.98 24.0  TRUE
2   9.14 21.6  TRUE
3   4.03 34.7  TRUE
4   2.94 33.4  TRUE
5   5.33 36.2  TRUE
6   5.21 28.7  TRUE
7  12.43 22.9  TRUE
8  19.15 27.1  TRUE
9  29.93 16.5  TRUE
10 17.10 18.9 FALSE
[1] "size of training data"

1.415 2.15

[1] "size of testing data"

1.91 2.15
```

1.3 (2) Cross validation for out-of-sample predictions

Next we split our training data into K folds. For every fold k , we train on the left over $K - 1$ folds and make predictions on fold k . We repeat this process for all M component models.

We will use a KNN neighbor regression, linear regression, polynomial regression, and regression tree to predict the median value of houses (medv).

```
[39]: require(FNN) # for the KNN model
      require(rpart) # for the TBR model

K = 10
```

```

training = training[,names(training)!='train']
trainingFolds = split(training,1:K)

dataSetOfPredictions = matrix()
for(k in 1:K){
  sprintf("Fold %d", k)

  outOfSample = trainingFolds[[k]]
  outOfSampleX = outOfSample[,names(outOfSample)!='medv']

  leftOver = setdiff(1:K,k)
  trainingSamples = do.call(rbind,trainingFolds[leftOver])

  trainY = trainingSamples[,names(trainingSamples)=='medv']
  trainX = trainingSamples[,names(trainingSamples)!='medv']

  # KNN model
  m1_predictions = knn.reg(train = as.matrix(trainingSamples$crim)
    ,test = as.matrix(outOfSampleX$crim)
    ,y=as.matrix(trainY),k=10)$pred

  #linear regression
  m2 = lm(medv~., data = trainingSamples)
  m2_predictions = predict(m2,outOfSampleX)

  #polynomial regression
  m3 = lm(medv~tax + age + rm + crim + I(tax^2) + I(age^2) + I(rm^2) +
→I(crim^2), data = trainingSamples)
  m3_predictions = predict(m3,outOfSampleX)

  #TBR
  m4 = rpart(medv ~ ., method="anova", data=trainingSamples)
  m4_predictions = predict(m4,outOfSampleX)

  # build data set of out-of-sample predictions
  if (k==1){
    allDataSetOfPredictions =
→cbind(m1_predictions,m2_predictions,m3_predictions,m4_predictions,trainY)
  } else {
    dataSetOfPredictions =
→cbind(m1_predictions,m2_predictions,m3_predictions,m4_predictions,trainY)
    allDataSetOfPredictions =
→rbind(allDataSetOfPredictions,dataSetOfPredictions,trainY)
  }
}

```

Warning message in split.default(x = seq_len(nrow(x)), f = f, drop = drop, ...):

```

â€”data length is not a multiple of split variableâ€”
Warning message in cbind(m1_predictions, m2_predictions, m3_predictions,
m4_predictions, :
â€”number of rows of result is not a multiple of vector length (arg 1)â€”
Warning message in cbind(m1_predictions, m2_predictions, m3_predictions,
m4_predictions, :
â€”number of rows of result is not a multiple of vector length (arg 1)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions, trainY):
â€”number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in cbind(m1_predictions, m2_predictions, m3_predictions,
m4_predictions, :
â€”number of rows of result is not a multiple of vector length (arg 1)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions, trainY):
â€”number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in cbind(m1_predictions, m2_predictions, m3_predictions,
m4_predictions, :
â€”number of rows of result is not a multiple of vector length (arg 1)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions, trainY):
â€”number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in cbind(m1_predictions, m2_predictions, m3_predictions,
m4_predictions, :
â€”number of rows of result is not a multiple of vector length (arg 1)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions, trainY):
â€”number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in cbind(m1_predictions, m2_predictions, m3_predictions,
m4_predictions, :
â€”number of rows of result is not a multiple of vector length (arg 1)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions, trainY):
â€”number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in cbind(m1_predictions, m2_predictions, m3_predictions,
m4_predictions, :
â€”number of rows of result is not a multiple of vector length (arg 1)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions, trainY):
â€”number of columns of result is not a multiple of vector length (arg 3)â€”
Warning message in cbind(m1_predictions, m2_predictions, m3_predictions,
m4_predictions, :
â€”number of rows of result is not a multiple of vector length (arg 1)â€”
Warning message in rbind(allDataSetOfPredictions, dataSetOfPredictions, trainY):

```

â€”Number of columns of result is not a multiple of vector length (arg 3)â€”

We now have a dataset that includes an out-of-sample for prediction, for all 4 models and for all observation in our training set.

```
[40]: print(head(allDataSetOfPredictions))
```

m1_predictions	m2_predictions	m3_predictions	m4_predictions	trainY
34.69	30.747363	24.90661	25.3775	21.6
26.15	20.473482	21.74700	17.5925	20.4
22.13	15.836745	19.24050	17.5925	13.9
27.18	23.911609	21.00925	21.0876	20.0
18.76	8.919021	19.31691	17.5925	19.7
24.14	22.014722	27.78422	27.4200	33.0

1.4 (3) Build Aggregator

The next step trains a model that aggregates the M models together by training on the data set of out-of-sample predictions. We can consider a linear regression model as our aggregator.

```
[43]: agg = lm(trainY~m1_predictions+m2_predictions+m3_predictions+m4_predictions,data_
      ↪= data.frame(allDataSetOfPredictions))
      print(summary(agg))
```

Call:

```
lm(formula = trainY ~ m1_predictions + m2_predictions + m3_predictions +
    m4_predictions, data = data.frame(allDataSetOfPredictions))
```

Residuals:

Min	1Q	Median	3Q	Max
-21.052	-5.629	-1.674	3.533	36.329

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	10.74016	0.60623	17.716	< 2e-16 ***
m1_predictions	0.15542	0.03025	5.137	2.93e-07 ***
m2_predictions	0.26294	0.03494	7.526	6.52e-14 ***
m3_predictions	-0.01661	0.03966	-0.419	0.675426
m4_predictions	0.12455	0.03536	3.522	0.000433 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.865 on 3739 degrees of freedom

Multiple R-squared: 0.1374, Adjusted R-squared: 0.1365

F-statistic: 148.9 on 4 and 3739 DF, p-value: < 2.2e-16

These estimates are for our training set. We still need to make predictions on our held out test set.

1.5 (4) make predictions on test set

1.5.1 (4.1) Train Models on whole training set

```
[50]: # KNN model
      m1_predictions = knn.reg(train = as.matrix(training$crim)
                              ,test = as.matrix(testing$crim)
                              ,y=as.matrix(training$medv),k=10)$pred

      #linear regression
      m2 = lm(medv~., data = training)
      m2_predictions = predict(m2,testing)

      #polynomial regression
      m3 = lm(medv~tax + age + rm + crim + I(tax^2) + I(age^2) + I(rm^2) +
      →I(crim^2), data = training)
      m3_predictions = predict(m3,testing)

      #TBR
      m4 = rpart(medv ~ ., method="anova", data=training)
      m4_predictions = predict(m4,testing)

      testSetPredictions =
      →cbind(m1_predictions,m2_predictions,m3_predictions,m4_predictions)
      AggregatorModelPredictionsOnTestSet = predict(agg,data.frame(testSetPredictions))

      print(AggregatorModelPredictionsOnTestSet)
```

10	12	19	22	34	40	42	44
20.69625	22.77449	20.35950	20.86065	19.54791	25.93482	25.08254	22.98151
50	56	59	62	63	70	74	83
20.20664	27.60575	22.62655	21.14669	23.37689	22.09355	22.00723	24.97694
87	90	112	113	114	115	116	123
22.64821	26.87657	24.46649	22.04813	21.27338	23.44954	21.11128	21.87311
124	126	134	143	152	156	168	169
20.06635	21.34642	20.75505	19.79120	22.74392	21.30573	24.01421	22.88878
177	181	184	195	207	213	215	219
24.02294	27.05077	25.40939	26.96737	22.36660	21.77531	18.44573	22.75528
223	229	235	242	243	244	249	251
27.40530	28.20853	26.48792	22.85583	22.82411	24.21718	21.74605	23.24770
261	265	270	272	273	279	283	301
28.25513	28.87083	23.95928	23.70026	24.20370	25.77376	30.13138	25.92524
308	317	320	321	327	329	330	332

25.80997	21.12504	23.66076	23.09643	23.48482	22.94679	24.30928	21.84448
338	346	352	357	359	362	364	382
22.56417	21.74137	23.43435	19.91901	22.79382	21.28101	21.95751	18.98937
386	389	392	399	412	416	437	446
16.37126	15.95998	20.87359	15.44610	18.78176	16.67029	18.10447	17.13334
447	451	456	458	462	466	469	470
20.59037	19.28533	20.35809	18.31982	21.57295	20.76852	18.66859	20.34247
476	478	482					
19.40579	17.35471	24.31944					

1.6 (5) Compare your aggregated model to the test set

```
[56]: testSetAndAggPredictions = data.frame('T' = testing$medv
                                           , 'P' = AggregatorModelPredictionsOnTestSet
                                           , 'M1' = m1_predictions
                                           , 'M2' = m2_predictions
                                           , 'M3' = m3_predictions
                                           , 'M4' = m4_predictions)

SSE_AGG = sum( (testSetAndAggPredictions$T - testSetAndAggPredictions$P)^2 )

SSE_M1 = sum( (testSetAndAggPredictions$T - testSetAndAggPredictions$M1)^2 )
SSE_M2 = sum( (testSetAndAggPredictions$T - testSetAndAggPredictions$M2)^2 )
SSE_M3 = sum( (testSetAndAggPredictions$T - testSetAndAggPredictions$M3)^2 )
SSE_M4 = sum( (testSetAndAggPredictions$T - testSetAndAggPredictions$M4)^2 )

print(SSE_AGG)
print(SSE_M1)
print(SSE_M2)
print(SSE_M3)
print(SSE_M4)
```

```
[1] 2154.833
[1] 3365.197
[1] 1110.28
[1] 964.3875
[1] 1307.686
```

```
[ ]:
```